Natural Language Processing (COM4513/6513) Week 5: Information Extraction

Prof. Jochen L. Leidner

⟨leidner@acm.org⟩

The University Of Sheffield.

University of Sheffield, Department of Computer Science

8 March 2018

Objectives

In this session, we aim to:

- get introduced to information extraction, its concepts and history
- learn about different approaches to IE

"Information Extraction" Defined

Information Extraction (IE) is

- the extraction of structured (relational) data from unstructured (= textual) sources
- a practically-motivated engineering discipline (models not necessarily inspired by nature)
- the use of natural language processing techniques to populate the slots of structured templates with appropriate fillers

IE was conceived as a shortcut to built useful systems when full text understanding based on syntactic parsing was beyond the state of the art.

IE Example

"Concepcion, 23 Aug 88 (Santiago Domestic Service) – Police sources have reported that unidentified individuals planted a bomb in front of a Mormon Church in Talcahuano District. The bomb, which exploded and caused property damage worth 50,000 pesos, was placed at a chapel of the Church of Jesus Christ of Latter-Day Saints located at No 3856 Gomez Carreno Street.
The shock wave destroyed a wall, the roof, and the windows of the church, but did not cause any injuries.
Carabineros bomb squad personnel immediately went to the location and discovered that the bomb was made of 50 grams of an-fo ammonium nitrate-fuel oil blasting agents and a slow fuse."

IE Example: Terrorist Attack Event

Concepcion, **23 Aug 88** "(Santiago Domestic Service) – Police sources have reported that **unidentified individuals** planted a **bomb** in front of **a Mormon Church** in **Talcahuano District**. The bomb, which exploded and caused **property damage worth 50,000 pesos**, was placed at a chapel of the Church of Jesus Christ of Latter-Day Saints located at No 3856 Gomez Carreno Street.
The shock wave destroyed a wall, the roof, and the windows of the church, but **did not cause any injuries**.
Carabineros bomb squad personnel immediately went to the location and discovered that the bomb was made of 50 grams of an-fo ammonium nitrate-fuel oil blasting agents and a slow fuse."
— MUC4 story TST4-MUC4-0001

The "Terrorist Attack" Template

Terrorist Attack	
Type of Attack:	
Perpretator:	
Target:	
Location:	
Time:	
Casualties:	
Injured:	
Material Damage:	

The "Terrorist Attack" Template

Terrorist Attack	
Type of Attack:	bomb (ATTACK>BOMBING)
Perpretator:	unidentified individuals (unknown)
Target:	a Mormon church
Location:	Talcahuano District (Chile>Talcahuano)
Time:	23 Aug 88 (1988-08-23 00:00:00)
Casualties:	_____ (0)
Injured:	did not cause any injuries (0)
Material Damage:	property worth 50,000 pesos (50000)

A Short History of Information Extraction

- 1981: NYU "'Linguistic String" project (N. Sager)
- 1982: DeJong's FRUMP system: 'sketchy scripts'
- Carnegie Group build JASPER IE system for Reuters (Andersen et al., 1986)
- 1987/1989: MUCK I+II: Naval operations messages
- 1991-1998 MUC 3-7: Message Understanding Contest
- 2000-2004: ACE: Automatic Content Extraction
  - from text spans to abstract entities
  - English, Chinese, Arabic
- 2010s: First neural approaches to IE

Information Extraction System Architecture

- **Preprocessor/Tokenizer**: split story into units and ultimately word tokens
- **Gazetteer**: lexical look-up of important (to your task) words/phrases
- **POS tagger**: tag/disambiguate words w.r.t. parts of speech
- **Chunk parser**: find basic noun and verb phrases
- **Named entity tagger**: identify and classify proper names
- **Relationship tagger**: find relations between entities
- **Event Template Analyzer**: populate fact/event templates
- The result is a **structured fact/event database**

IE System Architecture – Example

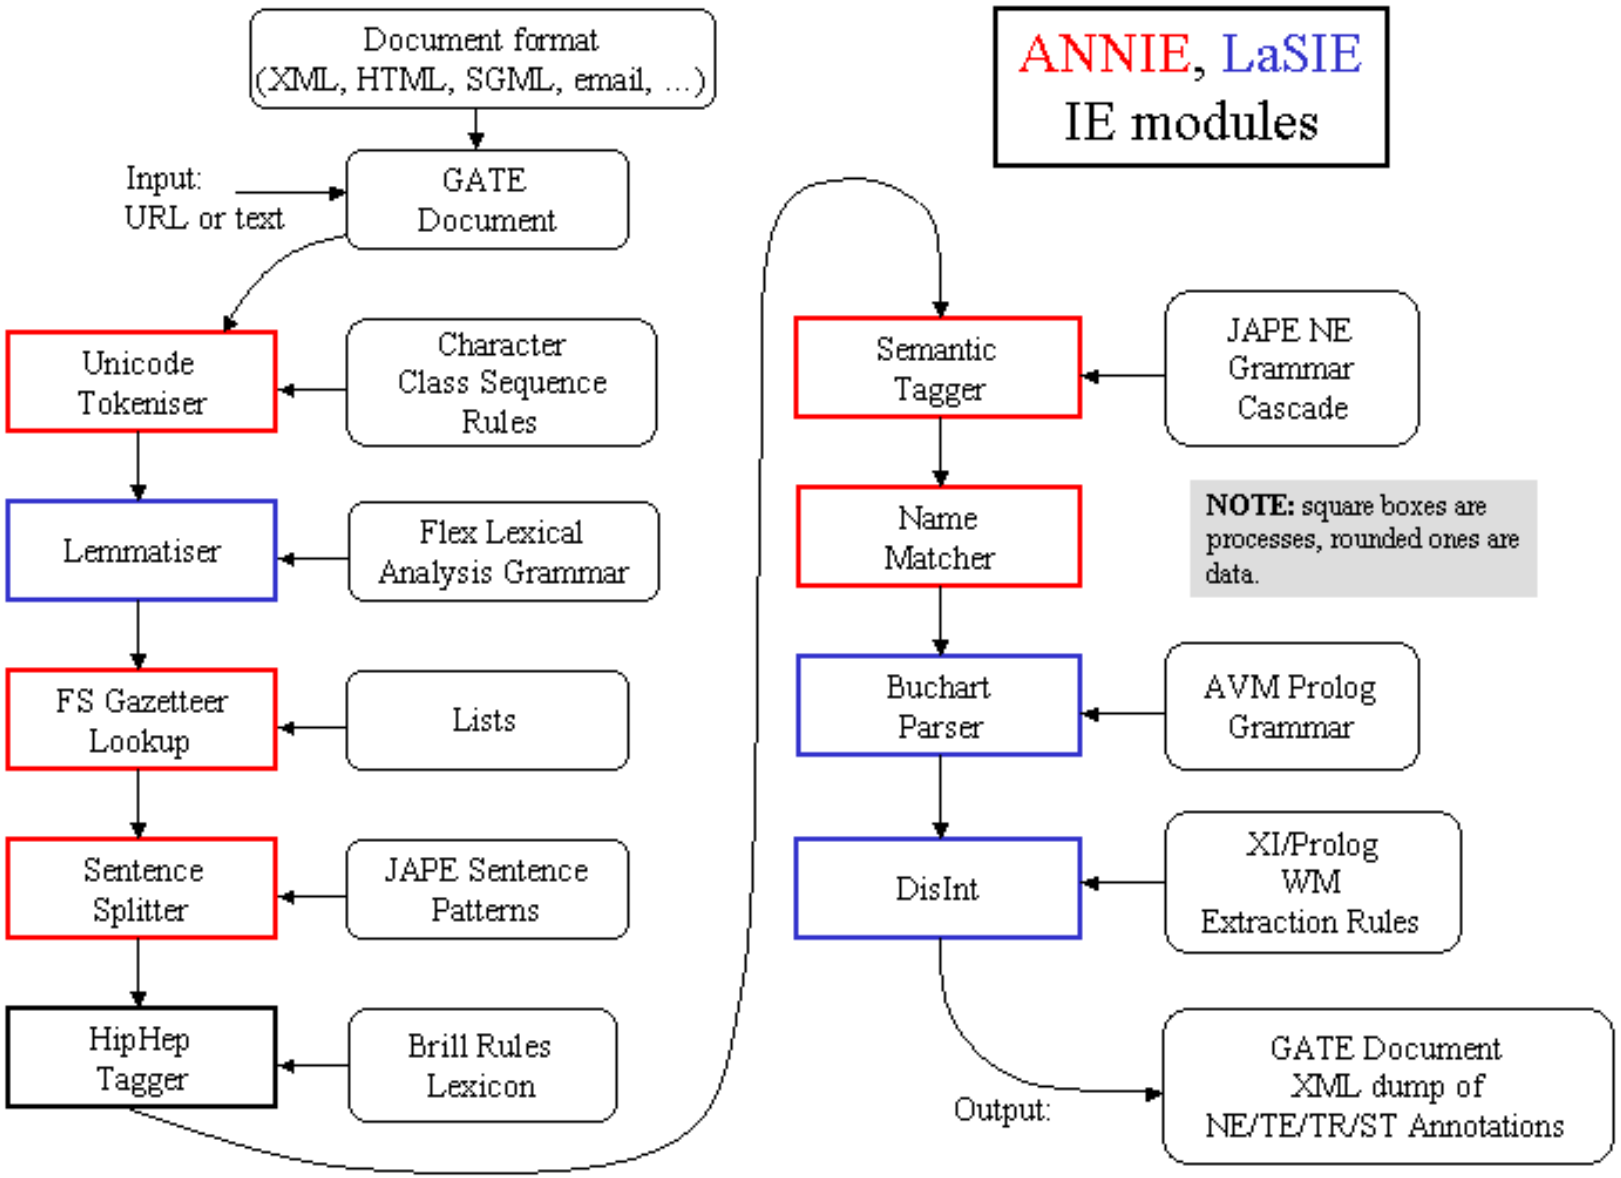

Fundamental Methods Used for IE

- **Rule-based**: human experts (computational linguists) manually write general linguistic rules and task-specific extraction rules.
  Trigger keywords, regular expressions, pattern/action, (cascaded) Finite State Transducers (FSTs)
- Supervised **Machine learning-based**: humans (domain experts) manually annotate text spans indicating entities, relations, facts etc. in a training corpus, and an expert (computational linguist) formulates a set of features; these get used to extract information if statistically correlated with classes of entities, relations etc. sought.
- Insight: shallow processing works well: SRI Tacitus $\rightarrow$ SRI FASTUS (Hobbs et al. 1992)

Example Patterns (FASTUS, Appelt et al., 1993)

killing of <HumanTarget>
<GovtOfficial> accused <PerpOrg>
bomb was placed by <Perp>
on <PhysicalTarget>
<Perp> attacked <HumanTarget>
<PhysicalTarget> with <Device>
<HumanTarget> was injured
<HumanTarget>'s body

From Entity over Relation to Scenario Template

- **Template Element Recognition (TE)**: extract information pertaining to organizations, persons and artifacts (NE tagging, parsing)
- **Relationship Extraction (RE)**: extract information about how individual entities stand in relationship to each other, drawing on a pre-defined inventory of relation types
- **Scenario Template Recognition (ST)**: extract pre-specified event information and relate the event to particular organizations, persons and/or artifacts (slot fillers)
- Information for filling a template often often spread across several sentences

Experimental Methodology

- Create a **gold data** set reference corpus with **ground truth**)
- Split gold data into three parts:
  - **development/training set**: used to study the data, train machine learning processes; can be inspected
  - **development test** ("devtest") set: cannot be inspected, cannot be used for training; repeatedly used to measure improvements of system quality by comparing system output with ground truth.
  - **test set**: cannot be inspected; only used once for final evaluation run at project end. Completely **unseen data** (to the system and developers).
- Gold data split:
  could be e.g. 80% train : 10% dev-test : 10% test

The GATE framework: From GUI to Guts

- GATE: General Architecture for Text Engineering - an open-source framework for *language engineering*
- Under development at the University of Sheffield for a long time (Cunningham et al., 2013)
- Current version: 8.4.2 (as of today) available from `http://gate.ac.uk`
- Java-based platform comprising GUI, APIs, and a workflow system
- GATE comes with pre-existing, contributed data (*Language Resources*) and components (*Processing Resources*)

Constructing IE Pipelines with GATE (1/2)

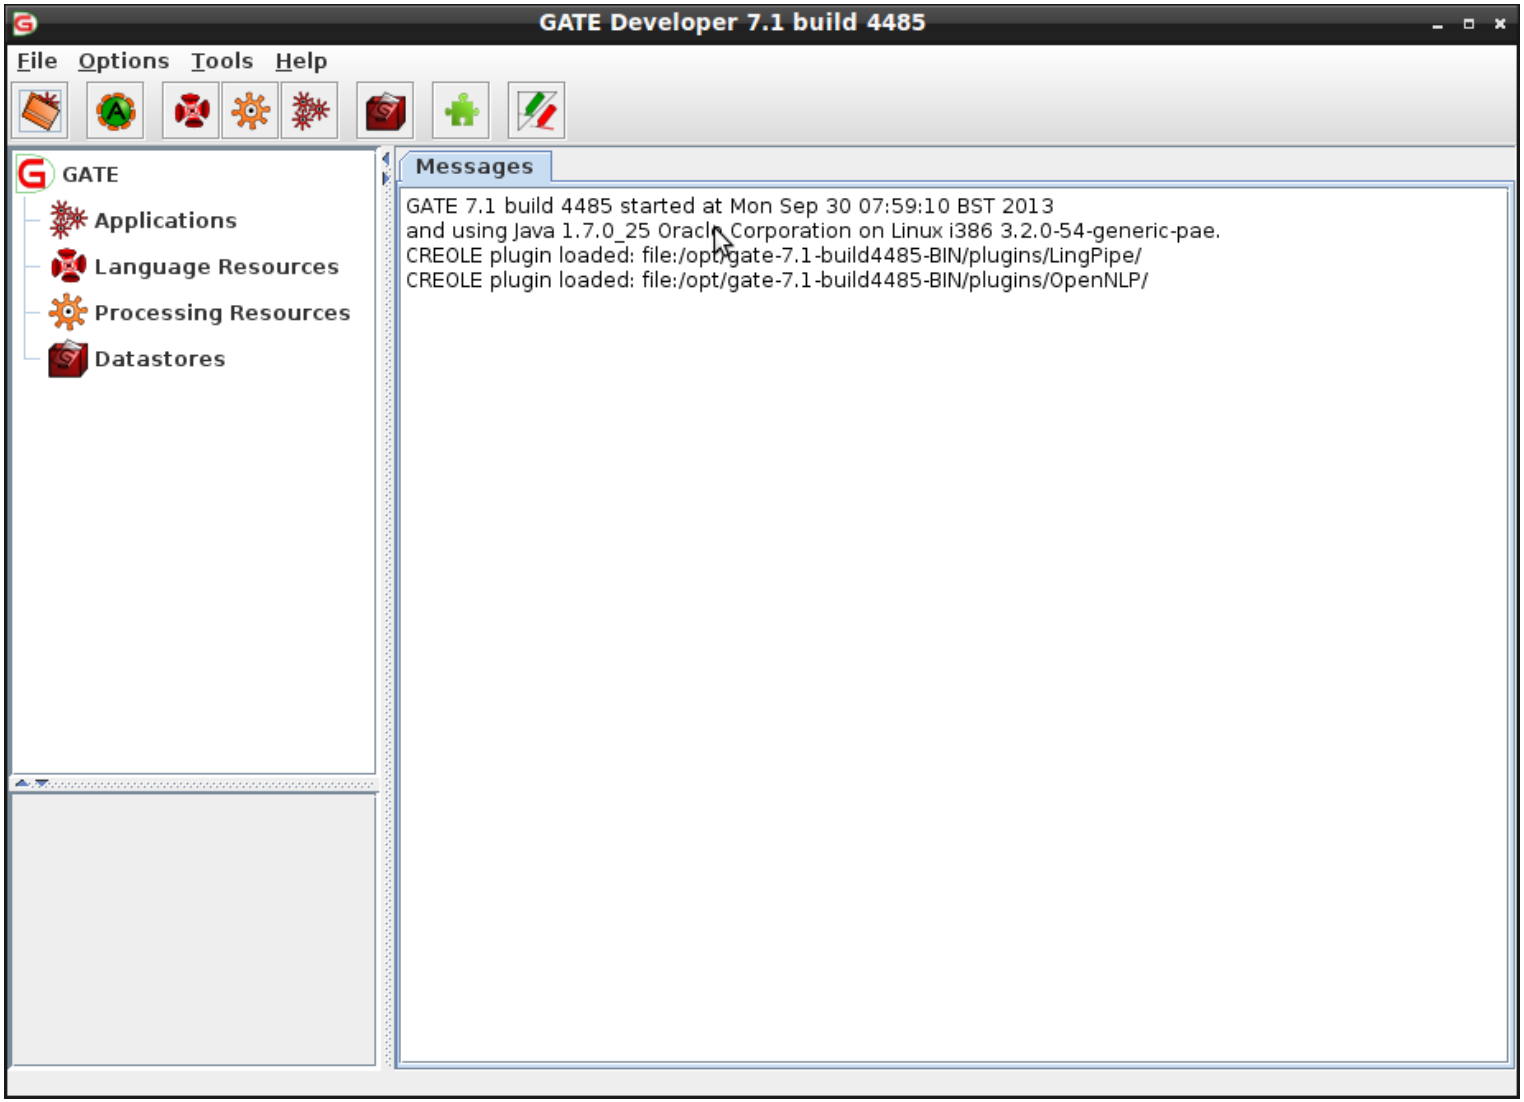

Constructing IE Pipelines with GATE (2/2)

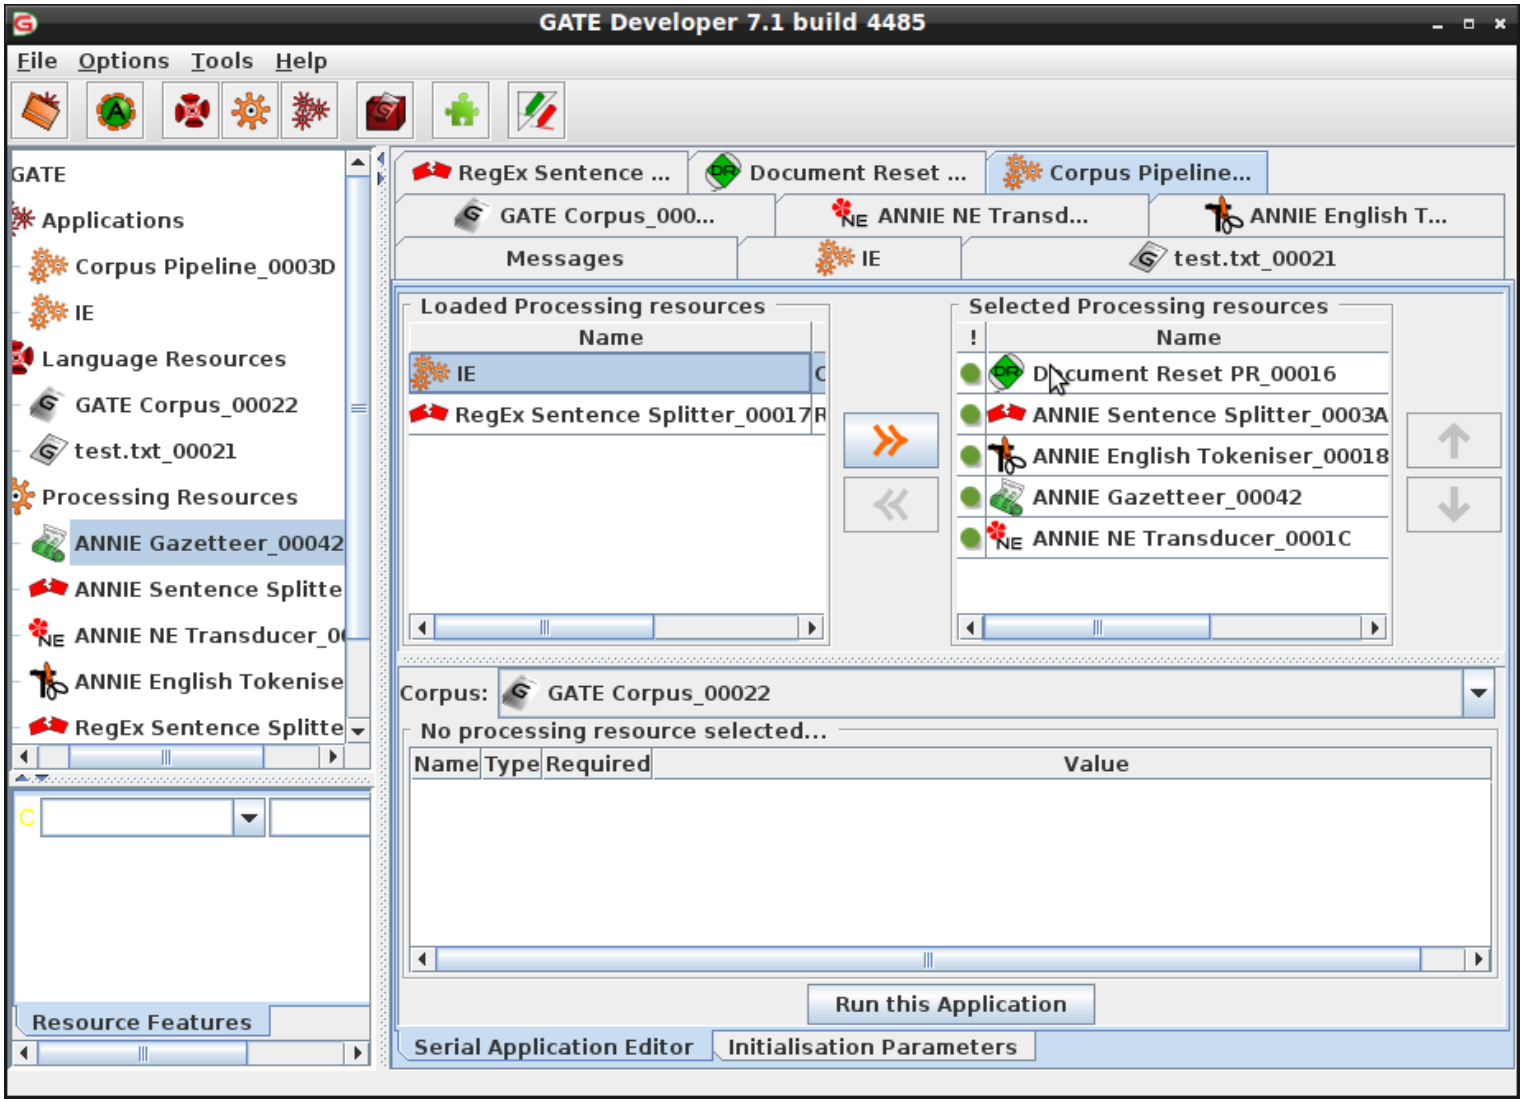

JAPE Introduction

- Part of the GATE platform
- Language to specify FSTs following the "patterns & actions" paradigm
- Sets of rules, broken down into processing phases
- Each rule tests matching conditions and typically adds annotations in the affirmative case

JAPE Rules – Regular Expressions over Annotations

Jape rules are organized into phases (cascades).
Each JAPE rule has three parts:

- **Header**: name of the rule
- **Left-Hand Side (LHS)**: pattern
- **Right-Hand Side (RHS)**: action – e.g. add annotations
- Structure:

```
Rule: TagUnknownName
(
 {Token.category == NNP}
):x
-->
 :x.Unknown = { kind = "PN", rule = TagUnknownName }
```

- See also: `http://gate.ac.uk/sale/tao/splitch8.html#chap:jape`

JAPE Rule Format – Example 1: URL Prefix

```
Phase: UrlPre
Input:  Token SpaceToken
/* important: specify input! */
Options: control = appelt

Rule: Urlpre
( (({Token.string == "http"} |
    {Token.string == "ftp"})
   {Token.string == ":"}
   {Token.string == "/"}
   {Token.string == "/"}
  ) |
  ({Token.string == "www"}
   {Token.string == "."}
  )
):urlpre
-->
 :urlpre.UrlPre = {rule = "UrlPre"}
```

Gazetteers in GATE

How to use lists of keywords/phrases:

- Create a `*.def` file listing all gazetteers (with their major/minor categories):

  ```
  cities.lst:location:city
  organizations.lst:organization
  surnames.lst:name:surname
  forenames.lst:name:forename
  ```

- List one key phrase, name or keyword per lin in each of these files, e.g. in `cities.lst`:

  ```
  Abu Dhabi
  Berlin
  Chicago
  Frankfurt
  ...
  ```

- Create a gazetteer Processing Resource in your application pipeline that references your `*.def` file.

Common Machine Learning Methods for Information Extration

- **Hidden Markov Models (HMMs)** (recall from Week 2)
- Conditional Random Fields (CRF)
- Support Vector Machine (SVM)
- Artificial Neural Networks (NNs) (will be covered in Week 7)

BIO Encoding for Labels Classifying Text Spans (CoNLL)

- Annotate word tokens with inside/outside of class information:
  `The_O Oracle_I-ORG CEO_O 's_O name_O`
  `is_O Larry_I-PER Ellison_I-PER ._O`
  (horizonal)
  or:
  `The O`
  `Oracle I-ORG`
  `CEO O`
  `...`
  (vertical)
- `B` tag used to demarcate adjacent `I` tags (otherwise, two adjacent "I" tokens part of the same text span could not be distinguished from the case of two adjacent, separate text spans)

Annotation of Gold Data with BRAT (http://nlp.brat.org)

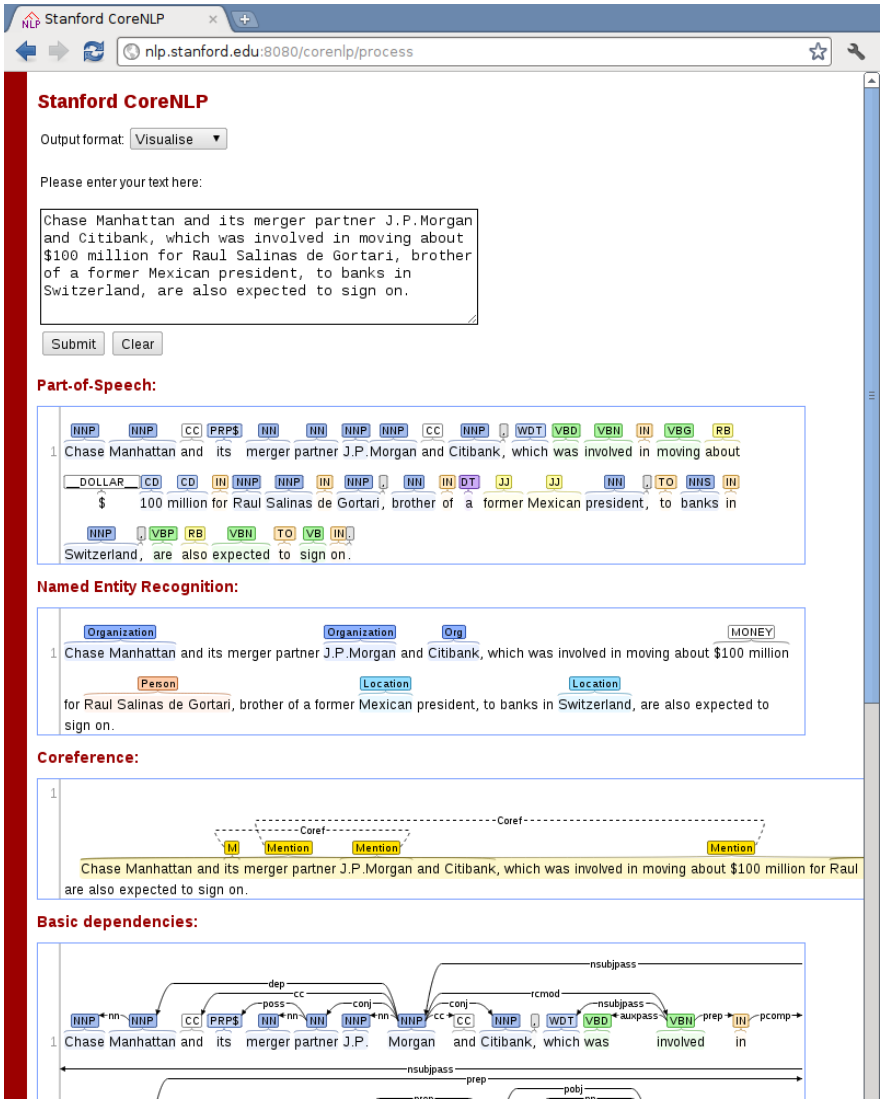

Feature Extraction – Spam Classification Example (1/2)

- **Feature**: a piece of evidence intended to help the classifier map the input to the right target class
- **Feature vector**: a vector $\vec{F}$, the components $F_j = \phi_j(d_i)$, of which are results applying a feature function to the data point d_i
- Example: "Spam versus Ham" email?
  - number of "!"s included in email body
  - length of the email in characters
  - does the word "cash" occur in the title or body?
- Example feature vectors:
  (2, 2392, no) $\mapsto$ `HAM` (genuine e-mail)
  (4, 520, yes) $\mapsto$ `SPAM`
  (1, 2392, no) $\mapsto$ `HAM`
  (0, 16337, no) $\mapsto$ `HAM`
  (1, 6i320, yes) $\mapsto$ `SPAM`

Feature Extraction – Features for Person NERC (2/2)

In English:

- Token begins with capital letter: `[A-Z][a-z]+`
- Token ends in characters `-man`
- Token to the left is from list Mr., Mrs., Miss, Dr., Prof., Sir, Lord, CEO, ...
- Token to the right is academic title, affiliation: BSc, Ph.D., M.A., FRS, M.P.
- Tokens to the right are `[,]? who`
- Token is from a list (gazetteer) of names

Training a Model

- Determine the model's parameters from data: induction
- Training corpus has counts, from which model probabilities are computed
- Smoothing: re-distribution of probability mass from seen to unseen events (to avoid zero probabilities, which make any product go zero)

HMMs for IE e.g. in Nymble (Bikel et al., 1997)

HMM Transitions:

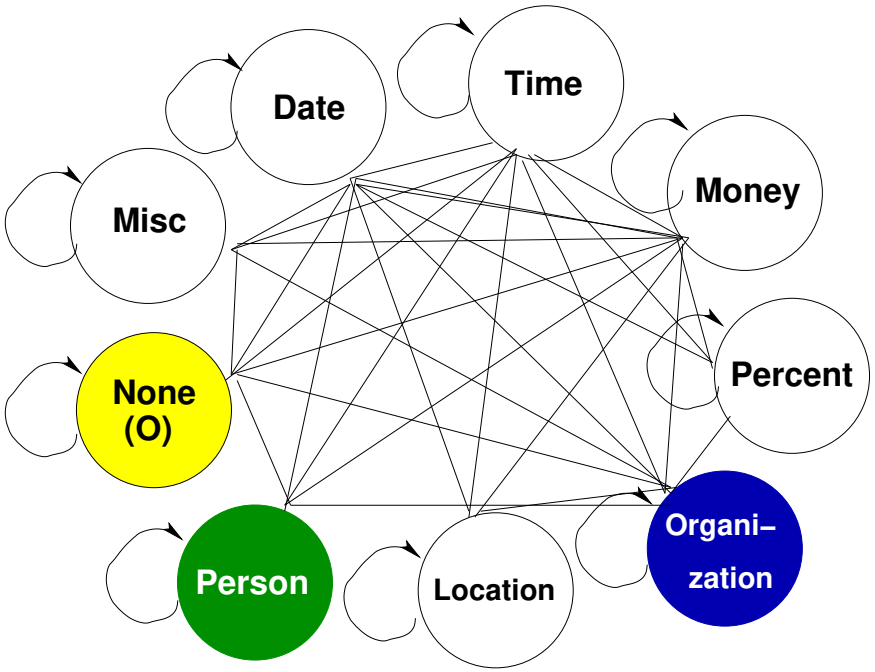

Not shown: initial tranistions/probabilities

Most likely state (label) sequence (Viterbi alg.):

Larry	Ellison	is	the	CEO	of	Oracle	Corporation	.
I–PER	I–PER	O	O	O	O	I–ORG	I–ORG	O

Features in Nymble (Bikel et al., 1997)

Word Feature	Example Text	Intuition
`twoDigitNum`	90	Two-digit year
`fourDigitNum`	1990	Four digit year
`containsDigitAndAlpha`	A8956-67	Product code
`containsDigitAndDash`	09-96	Date
`containsDigitAndSlash`	11/9/89	Date
`containsDigitAndComma`	23,000.00	Monetary amount
`containsDigitAndPeriod`	1.00	Monetary amount, percentage
`otherNum`	456789	Other number
`allCaps`	BBN	Organization
`capPeriod`	M.	Person name initial
`firstWord`	*first word of sentence*	No useful capitalization information
`initCap`	Sally	Capitalized word
`lowerCase`	can	Uncapitalized word
`other`	,	Punctuation marks, all other words

Quantifying System Quality

- Basis of system quality is a quantitative comparative evaluation
- When comparing predicted class (system) with actual class, there are 4 cases:
  - TP (true positive)
  - FP (false positive)
  - FN (false negative)
  - TN (true negative)
- **Confusion matrix**:

TP	FP
FN	TN

Accuracy, Precision and Recall

- **Sample size**: Number of data points N (e.g. tokens) to be classified
- **Accuracy** (A, Correctness): $A = \frac{TP+TN}{TP+TN+FP+FN} = \frac{C}{N}$
  is *not* a good measure to use
  $\rightarrow$artificially high values (if label bias)
- **Precision** (P, Positive Predictive Value): $P = \frac{TP}{TP+FP}$
  " fraction of retrieved instances that are relevant"
  # retrieved and correct over total retrieved (which proportion of responses were actually correct?)
- **Recall** (R, Sensitivity): $R = \frac{TP}{TP+FN}$
  "fraction of relevant instances that are retrieved"
  retrieved and correct over all correctly retrieved and missing (which proportion of correct responses was actually retrieved?)

F-Score

- Combined metric: **harmonic mean** of Precision and Recall
- Introduced by van Rijsbergen (1979)
- $F = (1 + \beta^2) \times \frac{PR}{\beta^2 P + R}$
- $F1 = F_{\beta=1} = \frac{2PR}{P+R}$ (equal weight of P and R)

Confusion Matrix (Source: Wikipedia)

- Most useful tool for performance evaluation

		predicted condition			
	total population	prediction positive	prediction negative	Prevalence $= \frac{\Sigma \text{ condition positive}}{\Sigma \text{ total population}}$	
true condition	condition positive	**True Positive (TP)**	**False Negative (FN)** (type II error)	True Positive Rate (TPR), Sensitivity, Recall, Probability of Detection $= \frac{\Sigma \text{ TP}}{\Sigma \text{ condition positive}}$	False Negative Rate (FNR), Miss Rate $= \frac{\Sigma \text{ FN}}{\Sigma \text{ condition positive}}$
	condition negative	**False Positive (FP)** (Type I error)	**True Negative (TN)**	False Positive Rate (FPR), Fall-out, Probability of False Alarm $= \frac{\Sigma \text{ FP}}{\Sigma \text{ condition negative}}$	True Negative Rate (TNR), Specificity (SPC) $= \frac{\Sigma \text{ TN}}{\Sigma \text{ condition negative}}$
	Accuracy $= \frac{\Sigma \text{ TP} + \Sigma \text{ TN}}{\Sigma \text{ total population}}$	Positive Predictive Value (PPV), Precision $= \frac{\Sigma \text{ TP}}{\Sigma \text{ prediction positive}}$	False Omission Rate (FOR) $= \frac{\Sigma \text{ FN}}{\Sigma \text{ prediction negative}}$	Positive Likelihood Ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$	Diagnostic Odds Ratio (DOR) $= \frac{\text{LR+}}{\text{LR−}}$
		False Discovery Rate (FDR) $= \frac{\Sigma \text{ FP}}{\Sigma \text{ prediction positive}}$	Negative Predictive Value (NPV) $= \frac{\Sigma \text{ TN}}{\Sigma \text{ prediction negative}}$	Negative Likelihood Ratio (LR−) $= \frac{\text{FNR}}{\text{TNR}}$	

Learning Curve: Do We Have Enough Gold Data?

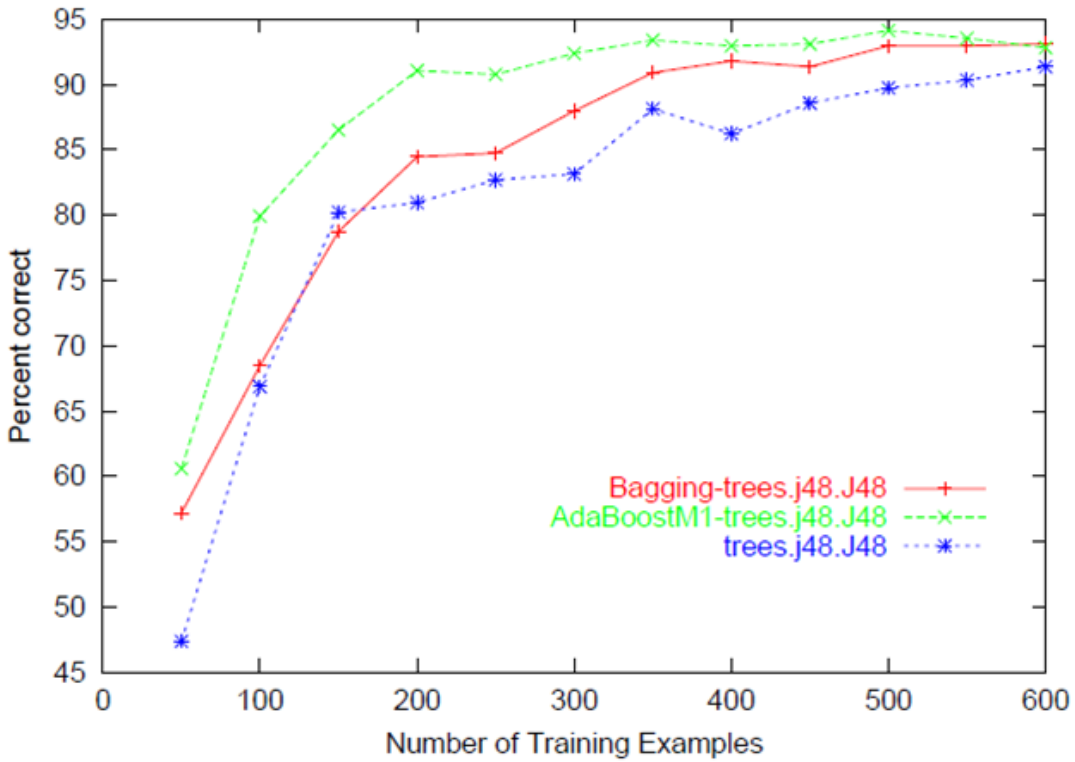

Real Applications of IE – Some Use Cases

There are many application domains for IE systems:

- Bio-medical IE applications (genes & proteins)
- Financial IE applications (mergers & acquisitions, CEO changes))
- Legal IE applications (judges & attorneys)
- Intelligence & Police IE applications (terrorists & crimes)
- e-Commerce IE applications (brands & products)

Event Extraction from News with Neural Models (Nugent et al., 2017) (1/2)

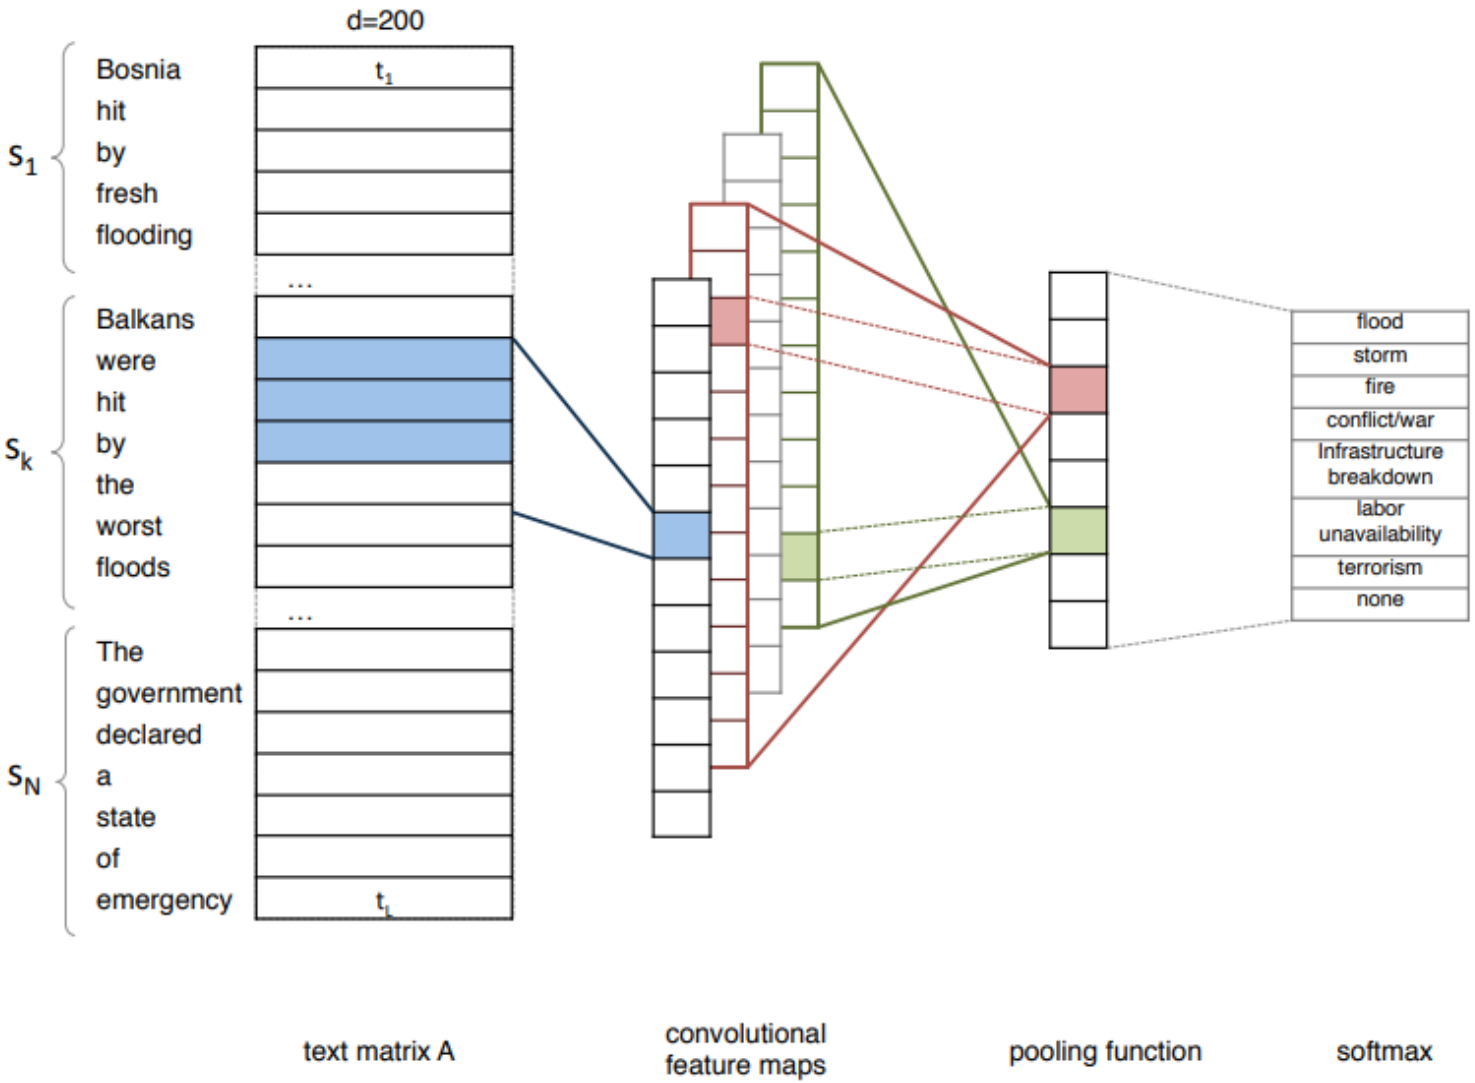

Event Extraction from News with Neural Models (Nugent et al., 2017) (2/2)

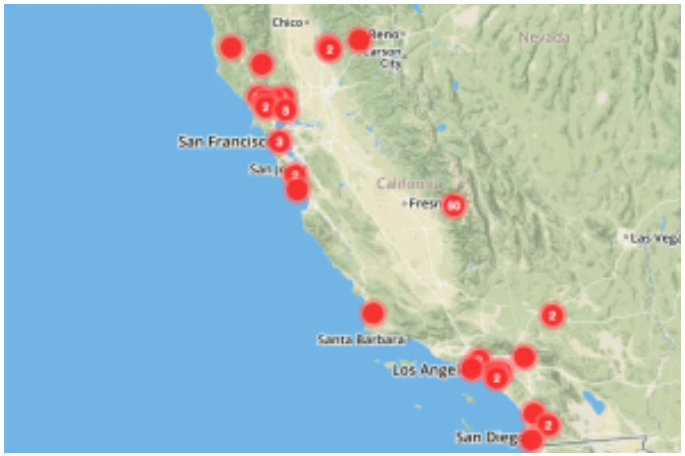

Wildfires in CA

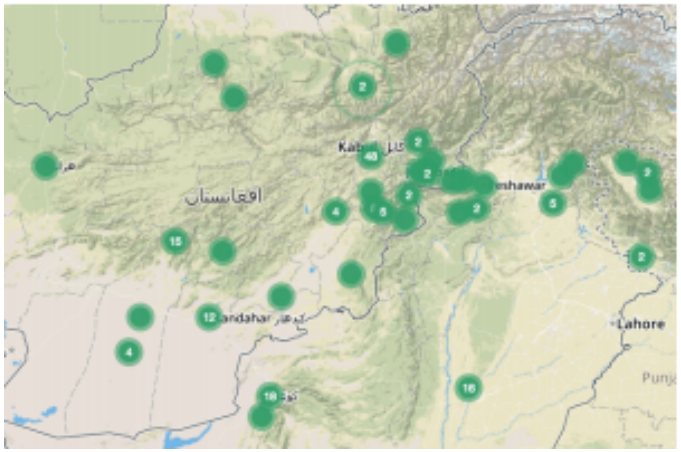

Conflicts in Afghanistan

Current Developments & Future Directions in IE

- Utilizing robust syntactic and semantic components where available
- From classic IE to *open IE*
- From small data sets to Web-scale data sets (bigger data and a simpler algorithm usually beats smaller data and a more sophisticaed algorithm!)
- Knowledge-rich methods are going to be combined with machine learning methods
- Use of word embeddings and deep (= multiple hidden layers) neural networks,

Summary

In this session, we learned:

- what information extraction is, and a bit of its history
- the difference between rule-based and machine learning approaches
- how named entities can be extracted from text

References (1/2)

- Bird/Klein/Loper (2009), *Natural Language Processing with Python*, Sebastopol, CA: O'Reilly
- Costantino and Coletti (2008), *Information Extraction in Finance*, Southampton: WIT Press
- Cunningham et al. (2013), Developing Language Processing Components with GATE Version 7 (a User Guide), University of Sheffield, [online] `http://gate.ac.uk/sale/tao/`, Chapter 2-3, 5-6, and 10.
- Hastie, Tibshirani and Friedman (2009), *The Elements of Statistical Learning (2nd ed.)*, New York, NY: Springer
- Jurafsky/Martin (2008), *Speech and Language Procesing (2nd ed.)*, Upper Saddle River, NJ: Prentice Hall
- Manning/Schütze (1999) *Statistical Natural Language Processing* MIT Press.
- Moens (2006), *Information Extraction: Algorithms and Prospects in a Retrieval Context* Dordrecht: Springer

References (2/2)

- Nugent & Leidner (2016), "Risk Mining: Company-Risk Identification from Unstructured Sources" *Proc. ICDM*
- Plachouras, Leidner & Garrow (2016) "Quantifying Self-Reported Adverse Drug Events on Twitter: Signal and Topic Analysis" *Ann. Meeting of the Social Media Soc.*
- Pustejovsky/Stubbs (2012), Natural Language Annotation for Machine Learning, Sebastopol, CA: O'Reilly, Chapters 5-6 and 8
- Zanasi (2005), *Text Mining and its Applications to Intelligence, CRM and Knowledge Management*, Southampton: WITPress

Backup Slides

Some Well-Known IE Systems

SAM	PAM	FRUMP	ATRANS	Proteus
PLUM	Tabula Rasa	LOLITA	Scrabble	NameTag
Alembic	Hasten	NLToolset	IE^2	TurboTag
SIFT	TASC	FACILE	AutoSlog	ANNIE
LasIE	Elie	NetOwl	Oki	BALIE
Diderot	JASPER	Tacitus	FASTUS	SCISOR
Circus	REES	Identifinder	Nymble	InfoXtract

Homework: Constructing IE Pipelines with GATE (1/2)

- File→New Corpus Pipeline, give name e.g. TestIE
- Processing Resources→Add Annie
- Processing Resources→Add Document Reset, ANNIE Sentence Splitter, ANNIE English Tokeniser, ANNIE Gazetteer, and ANNIE NE Transducer.
- Double click on TestIE and move the components Document Reset, ANNIE Sentence Splitter, ANNIE English Tokeniser, ANNIE Gazetteer, and ANNIE NE Transducer from the left list to the right using ">>"; ensure the order is exactly as given here from top to bottom!
- Language Resources→Add Corpus Document, click on sourceURL and select a text file to process

Homework: Constructing IE Pipelines with GATE (2/2)

- Language Resources→Add Corpus, then add document created above via drop-down menu to the empty corpus
- Double-click on TestIE Application, select corpus in drop-down and click Run this Application
- Double-click on your test document under Language Resources and select the annotations (PERSON, ORGANIZATION) you would like to view

Installing NLTK

```
$ sudo apt-get install python-nltk
$ python
>>> import nltk
>>> nltk.download()
```

Using NLTK – A KWIC Concordance

Let's extract a Keyword-in-Context concordance of the word "grail" in Monty Python and the Holy Grail:

```
from nltk.book import *
text6.concordance("grail")
```

```
ARTHUR : If you will not show us the grail , we shall take your
s required if the quest for the Holy grail were to be brought to
should separate , and search for the grail individually . [ clop
AD : You are the keepers of the Holy grail ? ZOOT : The what ? G
il ? ZOOT : The what ? GALAHAD : The grail . It is here . ZOOT :
ease ! In God ' s name , show me the grail ! ZOOT : Oh , you hav
rment me no longer . I have seen the grail ! PIGLET : There ' s
en the Grail ! PIGLET : There ' s no grail here . GALAHAD : I ha
are you going ? GALAHAD : I seek the grail ! I have seen it , he
 which , I have just remembered , is grail - shaped . It ' s not
blem . GALAHAD : It ' s not the real grail ? DINGO : Oh , wicked
```

Using NLTK – A POS Tagging Example

```
>>> import nltk
>>> sentence = """Although big data is a cool topic, Monday
morning lectures are never easy."""
>>> tokens = nltk.word_tokenize(sentence)
>>> tokens

>>> tagged = nltk.pos_tag(tokens)
['Although', 'big', 'data', 'is', 'a', 'cool', 'topic', ',
 'Monday', 'morning', 'lectures', 'are', 'never', 'easy',

>>> tagged
[('Although', 'IN'), ('big', 'JJ'), ('data', 'NNS'), ('is',
'VBZ'), ('a', 'DT'), ('cool', 'NN'), ('topic', 'NN'), (',',
','), ('Monday', 'NNP'), ('morning', 'NN'), ('lectures',
'NNS'), ('are', 'VBP'), ('never', 'RB'), ('easy', 'JJ'),
('.', '.')]
```

Using NLTK – the Built-in Named Entity Tagger

```
>>> nltk.ne_chunk(nltk.pos_tag(nltk.word_tokenize(
        "President Obama had to face the decision
         whether to visit Putin or not.")))
Tree('S', [('President', 'NNP'), Tree('PERSON', [('Obama',
'NNP')]), ('had', 'VBD'), ('to', 'TO'), ('face', 'VB'),
('the', 'DT'), ('decision', 'NN'), ('whether', 'IN'),
('to', 'TO'), ('visit', 'VB'), Tree('PERSON', [('Putin',
'NNP')]), ('or', 'CC'), ('not', 'RB'), ('.', '.')])
```

Using NLTK – A Regular Expression Tagger

```
>>> regexp_tagger = nltk.RegexpTagger(
...   [(r'^-?[0-9]+(.[0-9]+)?$', 'CD'),# cardinal numbers
...    (r'(The|the|A|a|An|an)$', 'AT'),# articles
...    (r'.*able$', 'JJ'),             # adjectives
...    (r'.*ness$', 'NN'),             # nouns formed from adjec
...    (r'.*ly$', 'RB'),               # adverbs
...    (r'.*s$', 'NNS'),               # plural nouns
...    (r'.*ing$', 'VBG'),             # gerunds
...    (r'.*ed$', 'VBD'),              # past tense verbs
...    (r'.*', 'NN')                   # nouns (default)
... ])
>>> regexp_tagger.tag(test_sent)
[('The', 'AT'), ('Fulton', 'NN'), ('County', 'NN'), ('Grand',
 'NN'), ('Jury', 'NN'), ('said', 'NN'), ('Friday', 'NN'),
 ('an', 'AT'), ('investigation', 'NN'), ('of', 'NN'),
 ("Atlanta's", 'NNS'), ('recent', 'NN'), ('primary', 'NN'),
 ('election', 'NN'), ('produced', 'VBD'), ('``', 'NN'),
 ('no', 'NN'), ('evidence', 'NN'), ....]
```

Using NLTK – Drawing Analysis Trees

```
from nltk.tree import *
from nltk.draw import tree
entities = nltk.ne_chunk(nltk.pos_tag(
    nltk.word_tokenize("Maria likes fast computers.")))
entities.draw()
```

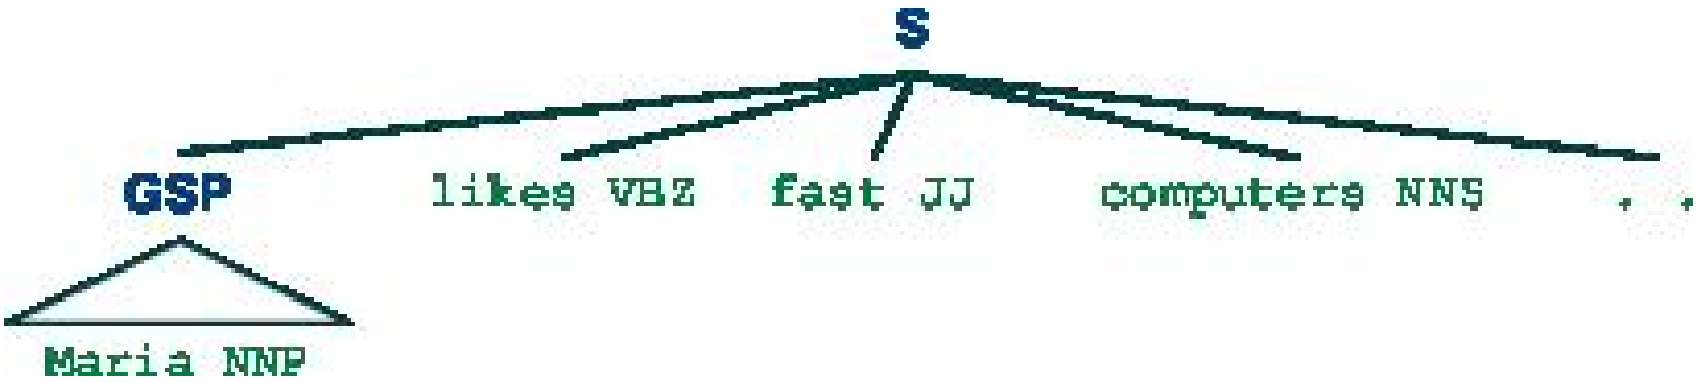

Sequence Tagging

- It is one thing to classify a data point in isolation
  Example (1): "`dinosaur`" $\mapsto$ noun
- It is quite another thing to classify, taking into account *context*
  Example (2a): "`(let's) walk (there)`" $\mapsto$ `walk/VVB`
  Example (2b): "`(a) walk (in the rain)`" $\mapsto$ `walk/NN`
- If "walk" has two possible readings, how can we compute the right one? $\rightarrow$**Ambiguity**
- Example (3): `I can can the can .` $\mapsto$ `I/PRP can/AUX can/VVB the/AT can/NN ./.`

Markov Model – Introductory Example (1/2)

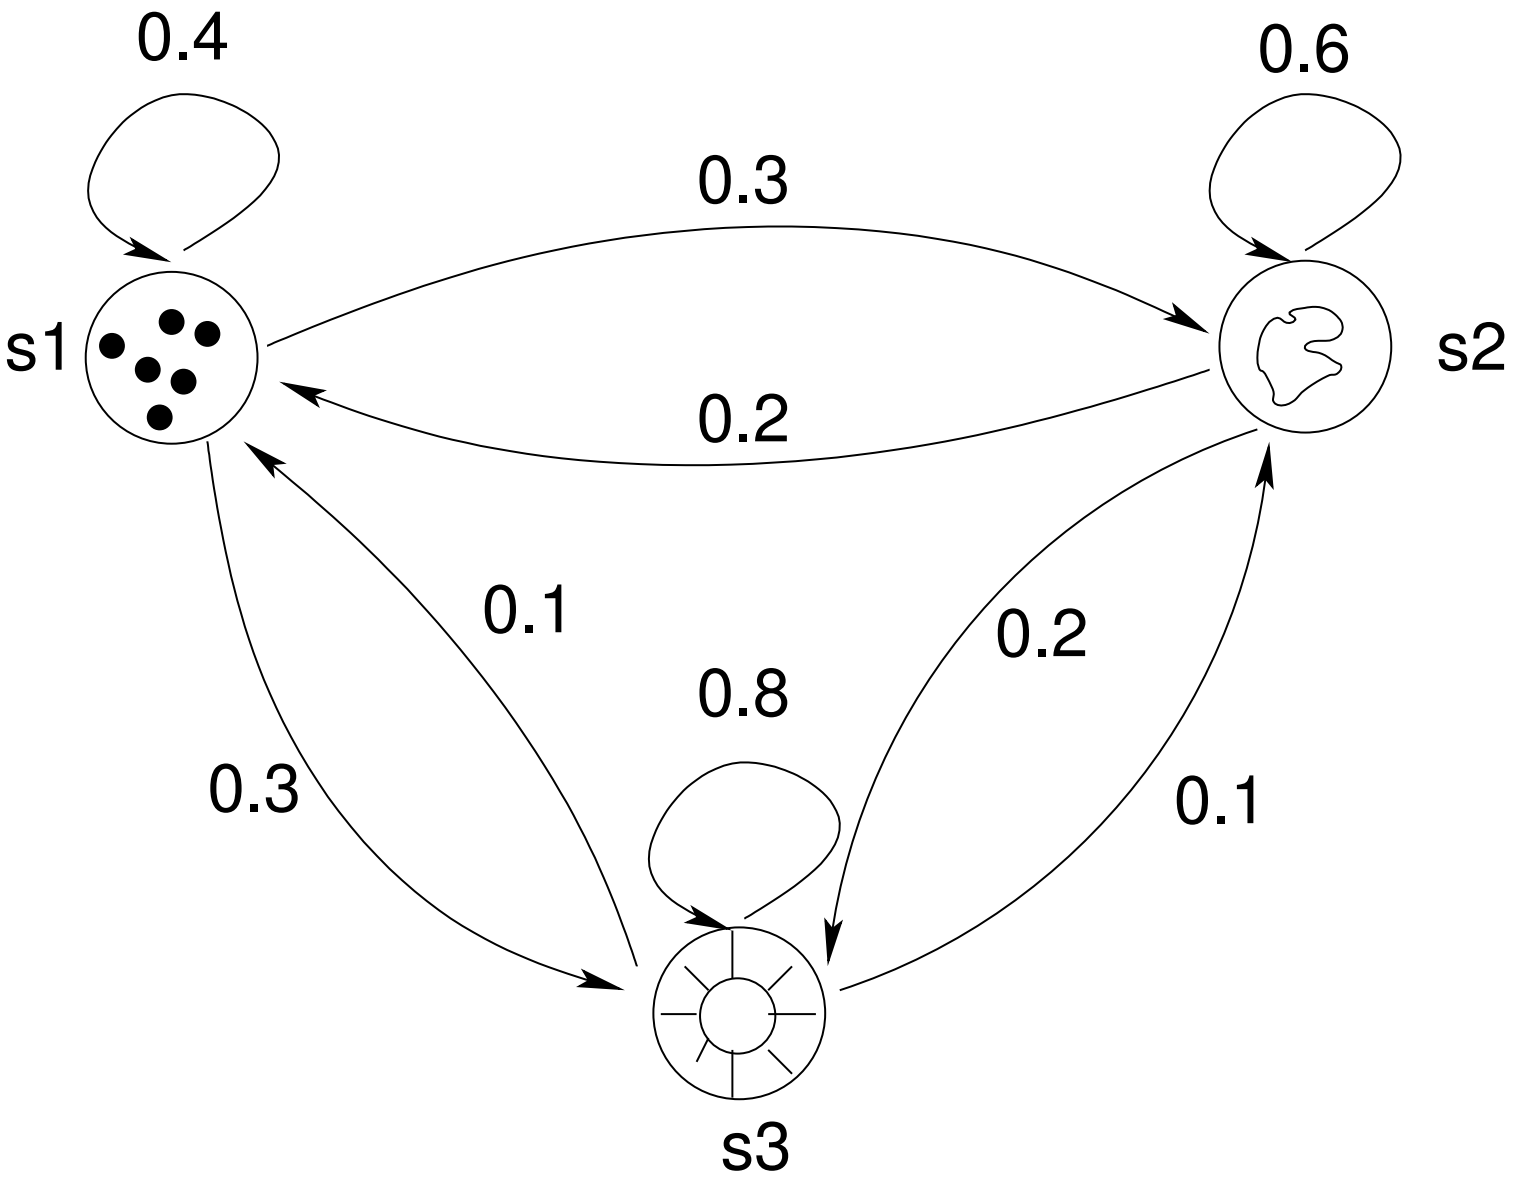

Markov Model – Introductory Example (2/2)

- Simple probabilistic finite-state model of the weather (**Markov Chain**)
- 3 States: s1: rainy, s2: cloudy, s3: sunny
- Transitions labelled with probabilities $a_{ij} \geq 0, \forall i, j \leq N$, $\sum_{j=1^N} a_{ij} = 1, \forall i$
- What is the probability of the observation sequence $O = (sunny, sunny, sunny, rainy, rainy, sunny, cloudy, sunny)$?

$$
\begin{aligned}
P(O|Model) &= P[s3, s3, s3, s1, s1, s3, s2, s3|Model] \\
&= P[3]P[3|3]^2 P[1|3]P[1|1]P[3|1]P[2|3]P[3|2] \\
&= \pi_3 {a_{33}}^2 a_{31} q_{11} a_{11} a_{13} a_{32} a_{23} \\
&= 1.0 \cdot 0.8^2 \cdot 0.1 \cdot 0.4 \cdot 0.3 \cdot 0.1 \cdot 0.2 \\
&= 1.536 \times 10^{-4}
\end{aligned}
$$

Hidden Markov Models (HMM)

A **Hidden Markov Model** $\lambda = (Q, \Sigma, A, B, \Pi)$ is formally defined as

- a set of N *hidden states* $Q = (q1, \ldots, q_N)$, $N = |Q|$
- a set Σ of M *observation symbols*, $M = |\Sigma|$
- a state-transition distribution (*transition probabilities*) $A = \{a_{ij} = P(q_{t+1} = j | q_t = i)\}, 1 \leq i, j \leq N$
- an observation symbol probability distribution (*emission probabilities*) $B = \{b_{ik} = b_o(o_k) = P(o_k | q_i)\}, 1 \leq k \leq M, o_k \in \Sigma$ (the probability that the output is o_k, given that the current state is q_i)
- an *initial state distribution* $\Pi = \{\pi_i = P(q_i | t = 1)\}, 1 \leq i \leq N$

Hidden Markov Models (HMM) – POS Tagging Example

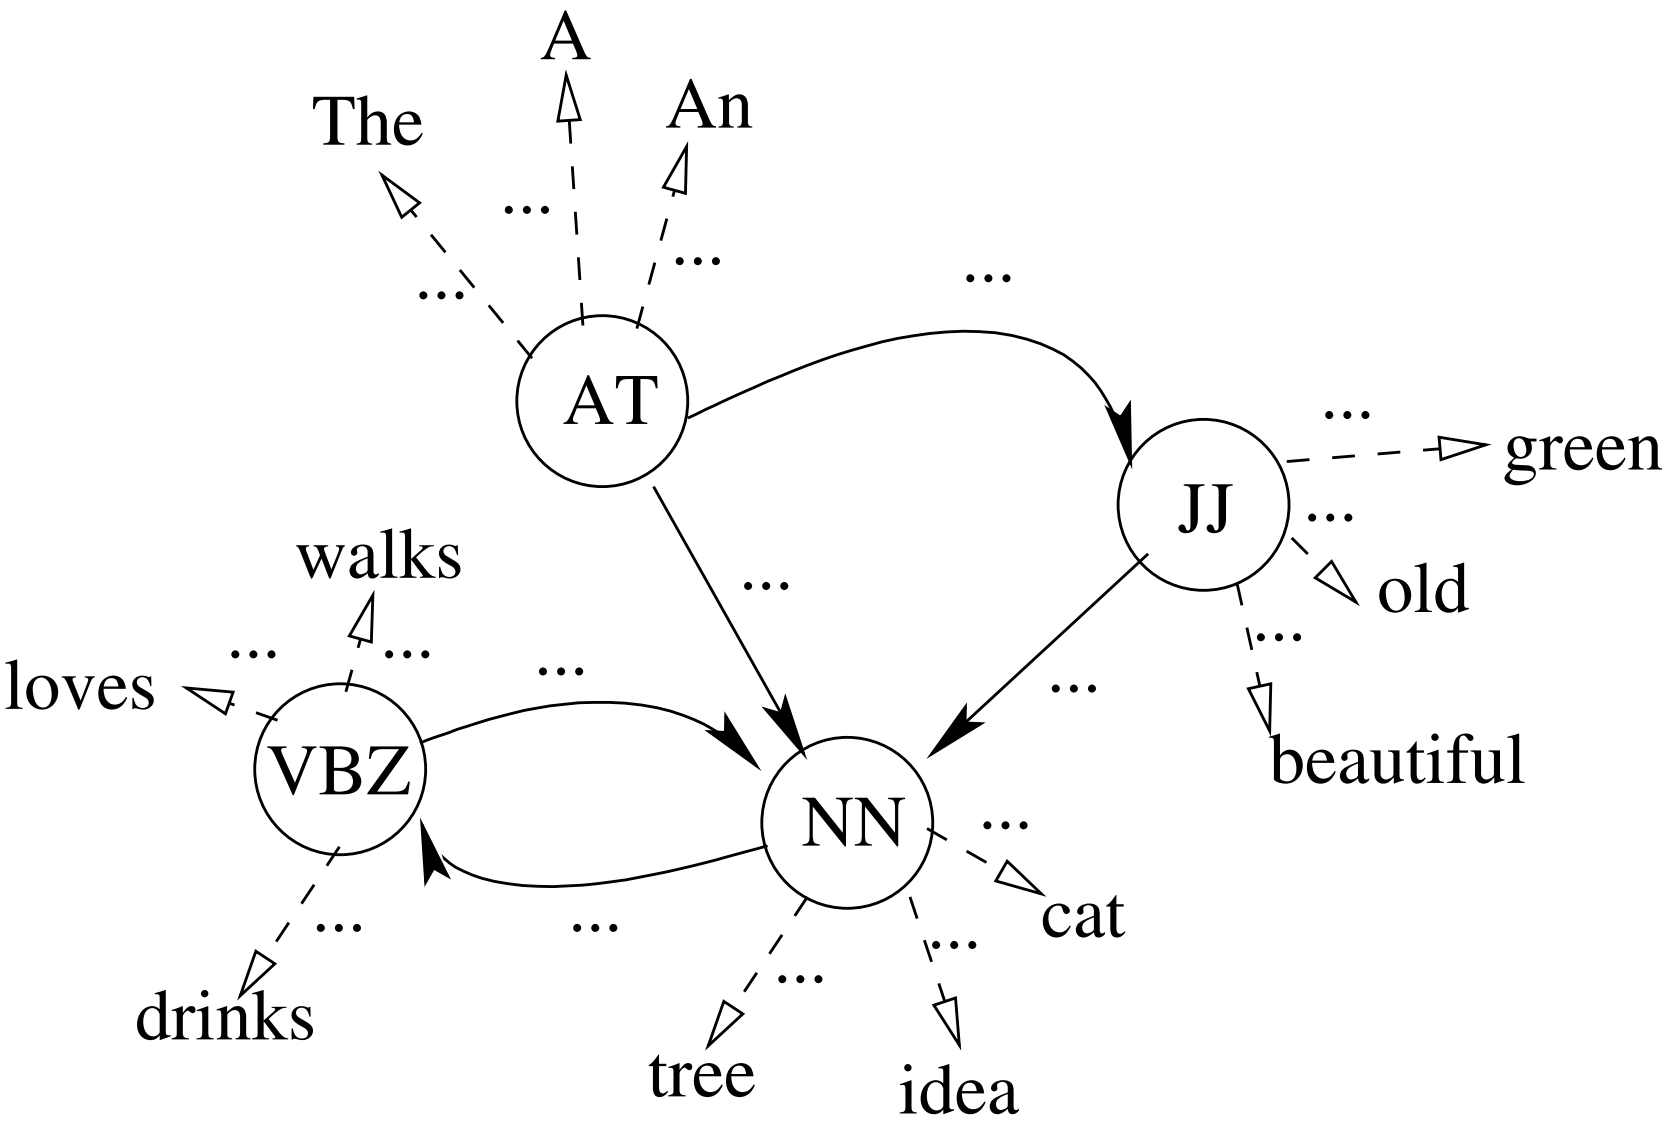

Three Questions to an HMM

- **Problem 1 (evaluation).** Given the observation sequence $O = (o_1, o_2, \ldots, o_T)$ and a model $\lambda = (Q, \Sigma, A, B, \Pi)$, how do we efficiently compute $P(O|\lambda)$, the probability of a observation sequence, given a model?
- **Problem 2 (decoding).** Given the observation sequence observation sequence $O = (o_1, o_2, \ldots, o_T)$ and the model λ, how can we find a corresponding state sequence $q^* = q_1, q_2, \ldots, q_T$ that most likely generated O ("optimally explains the observations")?
- **Problem 3 (learning).** How do we estimate the model parameters A, B and Π so as to maximize $P(O|\lambda)$?

The Viterbi algortihm (1/2) (Viterbi, 1967)

- Finding the most likely hidden state sequence, given an observation, requires at first glance exponenial runtime complexity ($O(N^n)$, as all possible states are valid hypotheses for each observation, so candidate states/readings multiply)
- Luckily, we can find a linear-runtime ($O(n)$) solution using dynamic programming (remembering partial solutions at each step) in a data structure called **trellis**.
- The Viterbi algorithm uses one pass from left to right looking at initial, transition and emission probabilities while storing a history of the locally "best" (most likely) state (back pointer).

Trellis Data Structure

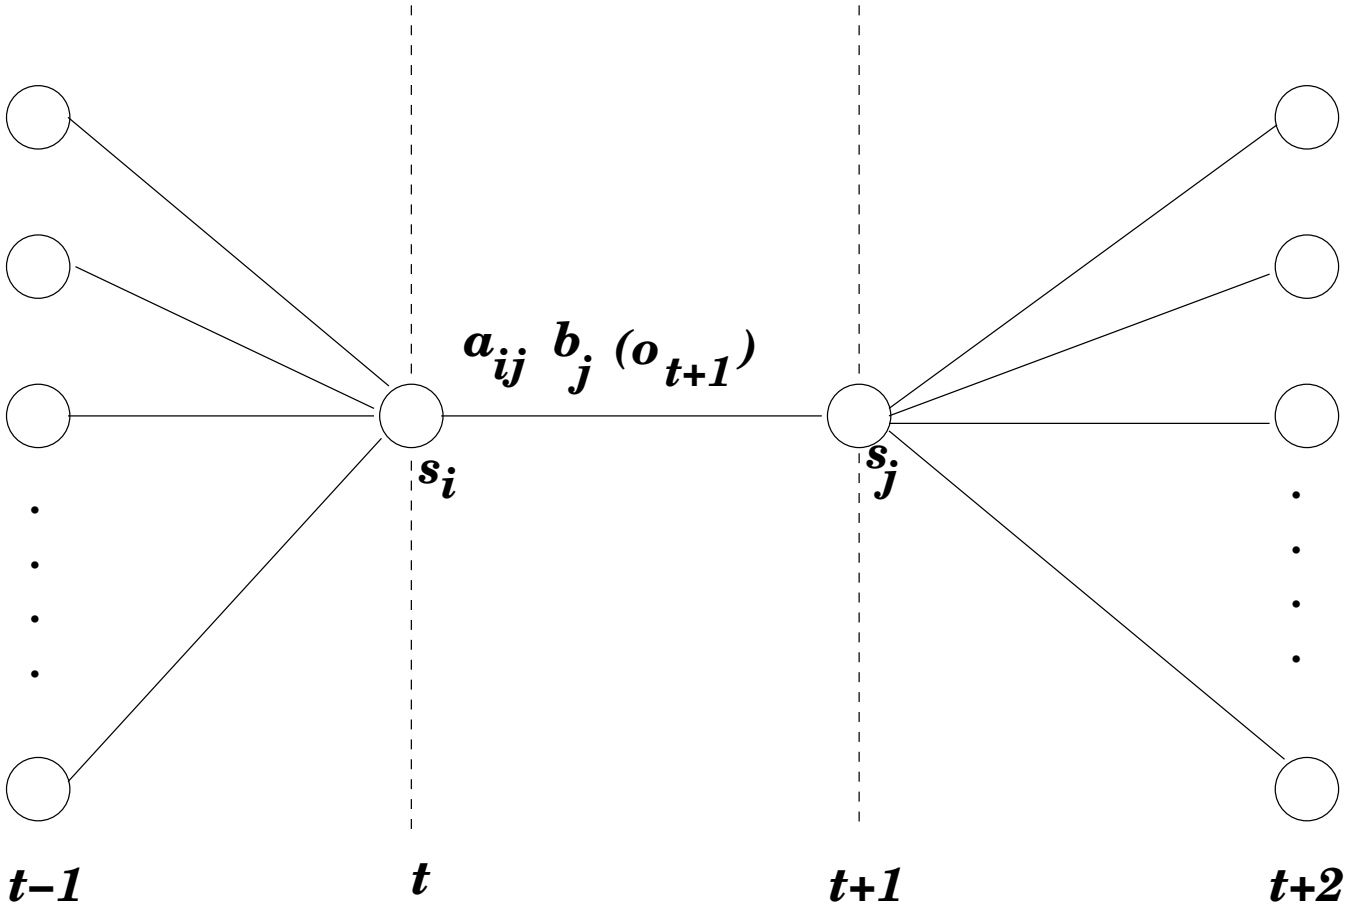

The Viterbi algortihm (2/2) (Viterbi, 1967)

Let $\delta_t(i) = \max_{q_1,\ldots q_{t-1}} P[q_1 q_2 \ldots q_{t-1}, q_t=i, o_1 o_2 \ldots o_t|\lambda]$.

1. Initialization. (Start with initial state probabilities)

$$\begin{aligned} \delta_1(i) &= \pi_i b_i(o_1), 1 \leq i \leq N \\ \phi_1(i) &= 0 \end{aligned}$$

2. Recursion. (main part, see previous slide)

$$\begin{aligned} \delta_t(j) &= \max_{1\leq i\leq N}[\delta_{t-1}(i)a_{ij}b_j(o_t)], 2 \leq t \leq T, 1 \leq j \leq N \\ \phi_t(j) &= \arg\max_{1\leq i\leq N}[\delta_{t-1}(i)a_{ij}], 2 \leq t \leq T, 1 \leq j \leq N \end{aligned}$$

3. Termination.

$$\begin{aligned} P^* &= \max_{1\leq i\leq N}[\delta_T(i)] \\ q_T^* &= \arg\max_{1\leq i\leq N}[\delta_T(i)] \end{aligned}$$

4. Path backtracking (finding the most likely hidden state sequence).

$$q_t^* = \phi_{t+1}(q_{t+1}^*), t = T-1, T-2, \ldots, 1$$

JRC European Media Monitor (EMM) (1/3) – Steinberger et al.

JRC European Media Monitor (EMM) (2/3)

EMM NewsBrief | EMM NewsExplorer | Name Search | Text Search

EMM NewsExplorer — Europe Media Monitor

People and Organisations

RSS feed for this entity

Daily News Analysis, across languages and over time

Main Menu

News Summary
About EMM NewsExplorer

News language and date

Language or country: en - English

Date: Nov 2013

Mo	Tu	We	Th	Fr	Sa	Su
				1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29	30	

Analysis over time

Timeline

Timeline [en] for 11/2013

50

Osama bin Laden

Information about this person was last updated on Saturday, November 16, 2013.

Names

Osama bin Laden (Eu,tr)
Bin Laden (da,tr)
bin Laden (da,tr)
Osama Bin Laden (da,sw)
بن لادن (ar,fa)
Oussama ben Laden (fr)
Ben Laden (de,ro)
Usama bin Ladin (de,tr)
Oussama Ben Laden (fr)
أسامة بن لادن (ar)
اسامة بن لادن (ar)
Осама бин Ладен (bg,sr)
Osama ben Laden (es,ro)
Osame bin Ladna (sl)
Bin Ladin (de,tr)
Usame Bin Ladin (tr)
Бин Ладен (bg,ru)
Usama bin Laden (de,sv)
Бен Ладен (bg,ru)
Osama bin-Laden (da,pt)
Usame bin Ladin (tr)
ابن لادن (ar)
ben Laden (fr,ro)

Key Titles and Phrases

al Qaeda leader (en - 284)
al-Qaeda leader (en - 204)
líder de Al Qaeda (es,pt - 126)
Al Qaeda leader (en - 94)
chef (da,fr - 546)
terrorchef (de - 101)
Al-Qaeda leader (en - 81)
terroristenführer (de - 75)
líder da Al Qaeda (pt - 62)
leader (en,it - 417)
terroristenleider (nl - 60)
führer (de - 103)
leider (nl - 157)
Al-Qaeda chief (en - 39)
al Qaeda chief (en - 36)
al-Qaeda chief (en - 30)
terroristenführers (de - 35)
líder da Al-Qaeda (pt - 32)
Al-Qaida leader (en - 30)
anführer (de - 63)
líder (es,pt - 103)
terroristenchef (de - 27)
born (en - 72)

External resources

Image obtained automatically from Wikipedia

Read Wikipedia entry

Explore Relations

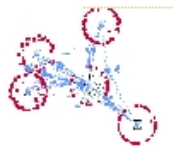

Related People

Barack Obama (447)
John Kerry (198)
Ayman al-Zawahiri (138)
George W. Bush (122)
Hillary Rodham Clinton (105)
Bashar Assad (101)
Kathryn Bigelow (78)
Hamid Karzai (75)
Chuck Hagel (75)
John McCain (74)
Nawaz Sharif (70)
Bill Clinton (64)
François Hollande (61)
Jay Carney (60)
David Cameron (60)
Vladimir Putin (59)
Steven Spielberg (58)
Ben Affleck (57)
Bradley Manning (56)
Edward Snowden (56)
Mohammed Morsi (55)
Daniel Day-Lewis (54)
Joseph Biden (53)
Leon Panetta (52)
Julian Assange (50)
Susan Rice (50)

Latest Clusters - English

[fr] [ar] [bg] [ru] [pt] [de] [hu] [da] [nl] [es] [no] [sv] [it] [tr] [ro] [sw] [et] [sl] [fa]

On the rise: Polio cases pass 2012 total
tribune 13-NOV-13

Blasts at Indian political rally kill 5
cnn 27-OCT-13

PM optimistic about investments in Pakistan

Afghan president criticises timing of death of Pakistan Taliban leader
telegraph 04-NOV-13

Sartaj Aziz says US understands Pakistan viewpoint on drone attacks
TheFrontierPost 25-OCT-13

President Obama welcomes PM Sharif at White House as

JRC European Media Monitor (EMM) (3/3)

Zoom:

Pan:

Search for a name:

Search results:

Click on a name to add it to the relation network:

These persons or organisations are already shown in the graphic:

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (1/6)

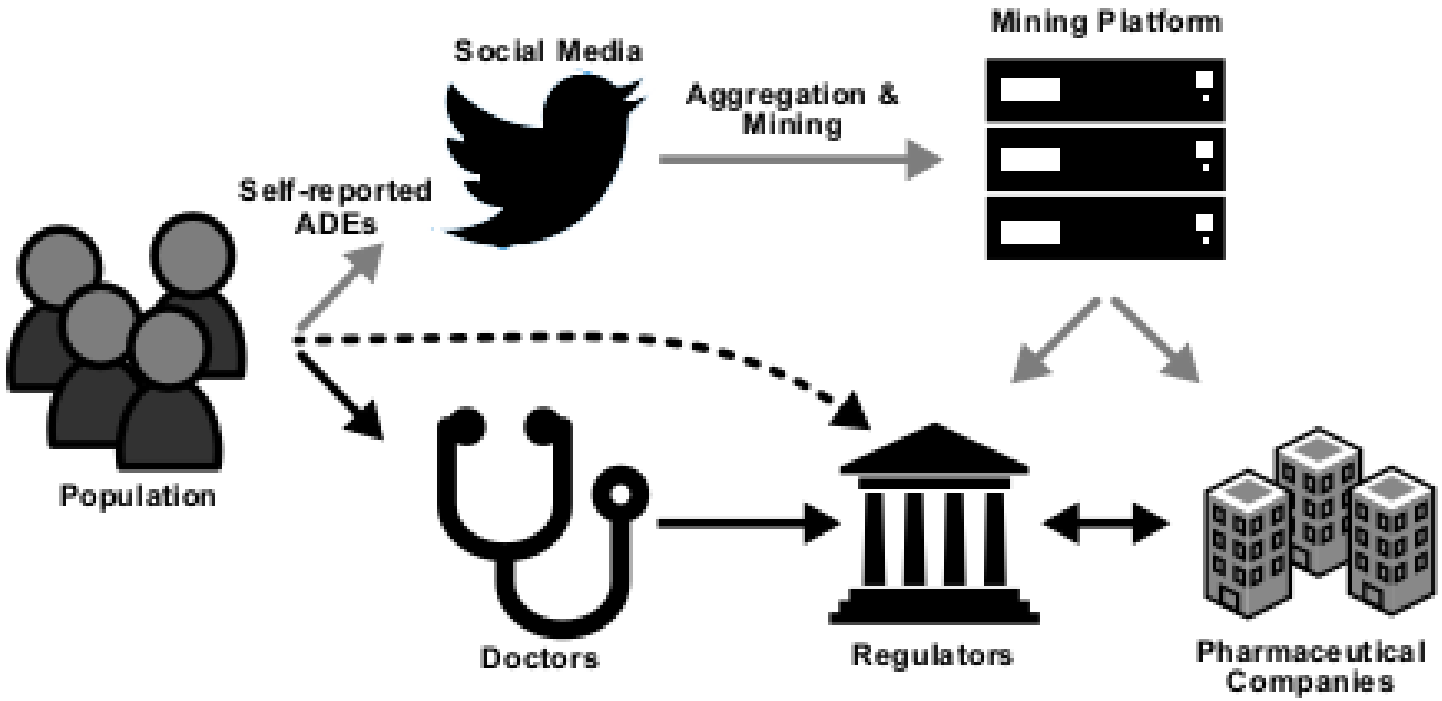

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (2/6)

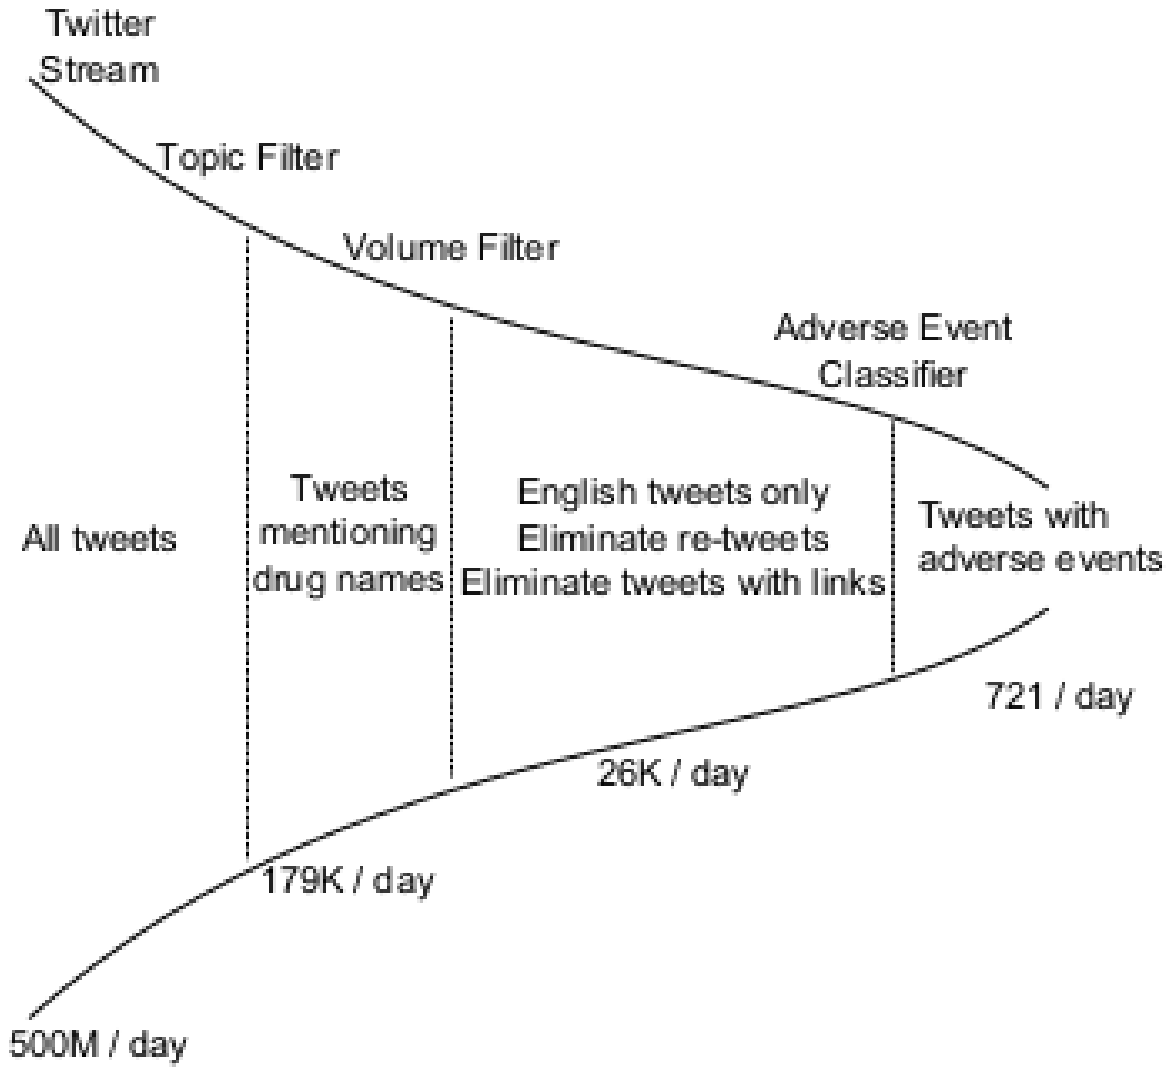

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (3/6)

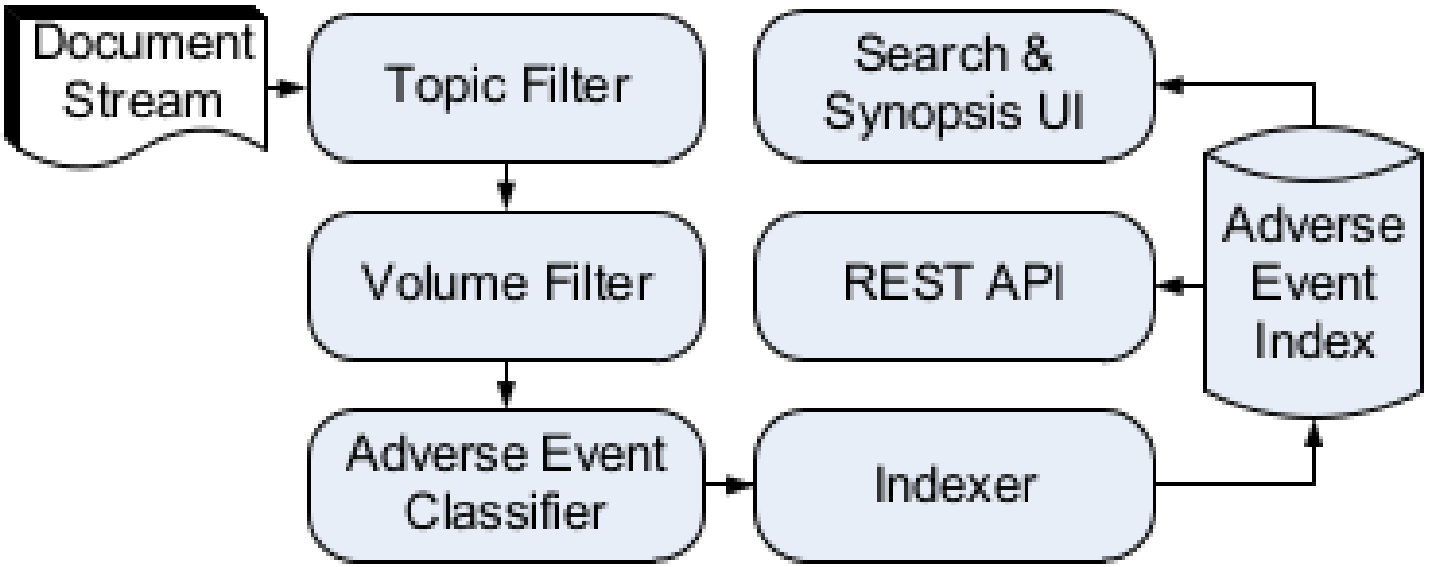

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (4/6)

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (5/6)

	Tweet text
TP	wellbutrin doesnt even work for me it just makes me really anxious idk why im still taking it
TP	Took some ibuprofen that has me so drowsy
TP	Insomnia and heart palpitations due to prednisone. Takteng mga side effects 'to. /wrist
TP	This Vicodin is making me feel like a tweaker because I'm so itchy!
TP	guys i took an ibuprofen 2 days ago and i still have heart palpitations
TN	I once gave a treatment to some patients with glutathione which I knew later clinic use it. Most of them have nausea for 5 mins after inj.
FN	I seriously never have any energy thanks accutane lol @probs_accutane all I want to do is sleep
FP	My mouth taste like 1200 Mg's of ibuprofen yet my head still hurts and I'm still feeling dizzy. Wtf
FP	@ash_hein they gave me Tylenol 3s & yea kinda my mouth still hurts a little & I'm still swollen

Pharmacology IE – Medical Drug Use Monitoring (Plachouras et al., 2016) (6/6)

			Development			Test			
Model	C	w^+	P	R	F1	P	R	F1	
Baseline			0.269	0.683	0.386	0.267	0.705	0.387	
BIN_NGRAM1,2	0.050	8	0.636	0.504	0.562	0.560	0.540	0.549	$*$
ALL FEATURES	0.025	9	0.573	0.590	0.582	0.550	0.669	**0.604**	$*\dagger$

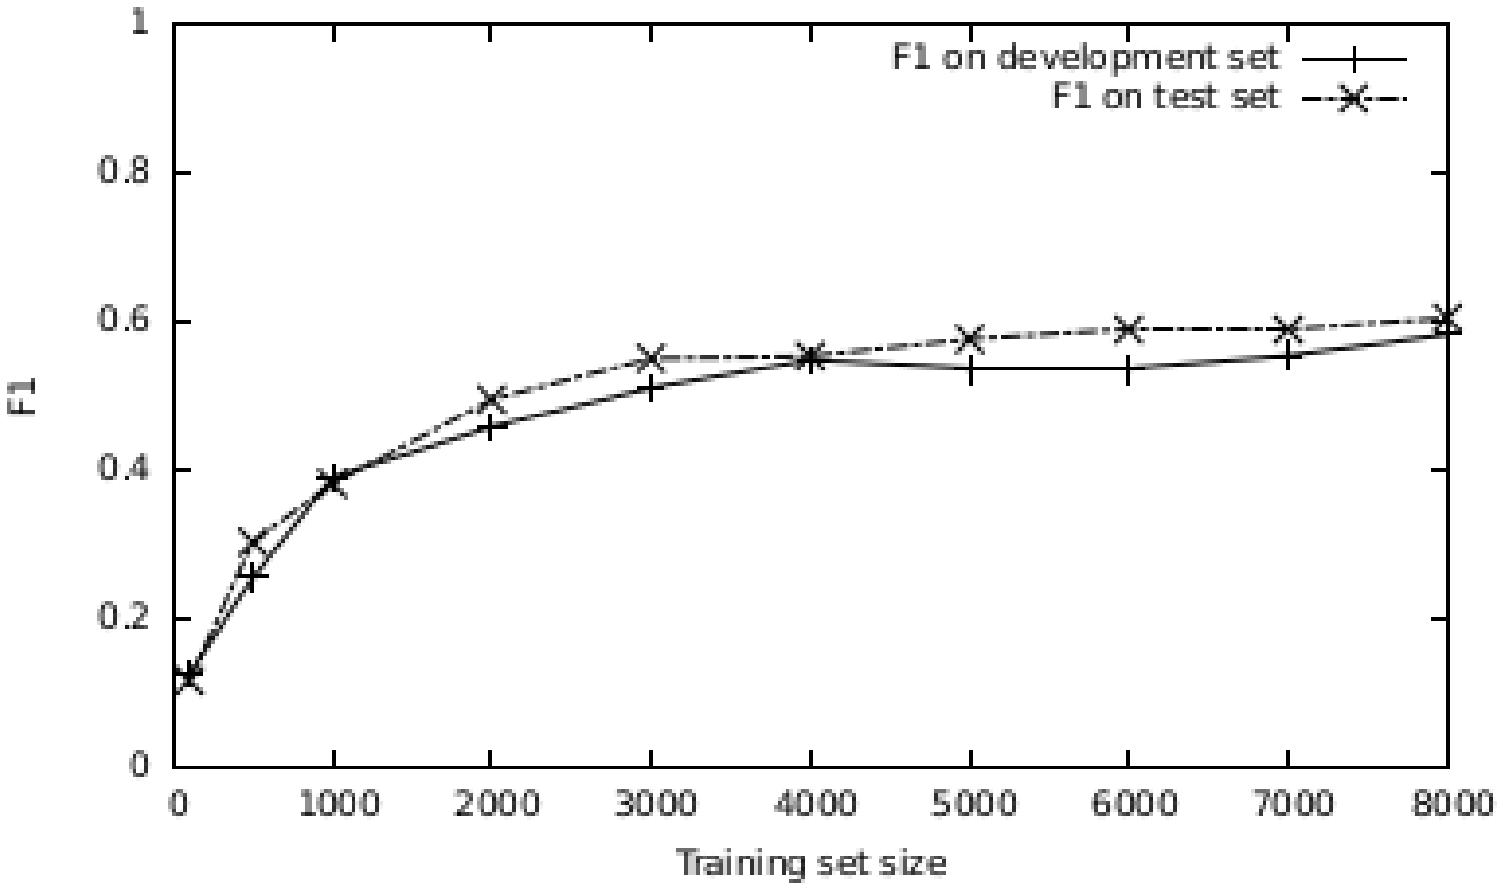

Where to Get Linguistic Gold Data From

- Search the Linguistic Data Consortium (LDC) catalog (`http://catalog.ldc.upenn.edu`)
- Search ELRA catalog (`http://catalog.elra.info`, `http://www.hlt-evaluation.org`)
- Browse DFKI's LT World `http://www.lt-world.org`
- Browse MetaNet (`http://www.meta-net.eu`)
- Browse the ACL Anthology (`http://acl.ldc.upenn.edu`), ACM Digital Library (`http://dl.acm.org`) and IEEExplore (`http://ieeexplore.ieee.org`)
  - some papers on a topic may contain associated data
  - some conferences, notably the bi-annual LREC, specialize on presenting new linguistic resources
- Contact expert researchers
- Recruit linguists/domain experts
- Crowdsourcing